

A Constrained Auditing Pipeline for Predictive HR Analytics: Explainability, Fairness Diagnostics, and Structural Limitations Demonstrated on a Synthetic Benchmark

Shaurya Gupta

Abstract

Predictive machine learning for employee attrition is routinely evaluated without principled constraints against causal overreach, calibration failure, or the amplification of small-sample variance risk. This paper proposes a heavily constrained auditing pipeline architecture — a normative methodological framework that structurally enforces explainability requirements, fairness diagnostics, and calibration checks prior to any deployment consideration in organizational settings. The framework is demonstrated on the IBM HR Analytics Employee Attrition dataset (1,470 observations, 35 variables), a synthetic, pedagogically-designed benchmark created for instructional purposes (IBM Data Scientists, 2015). The dataset is used as a fully controlled evaluation substrate for stress-testing the auditing mechanics, not as a source of novel organizational findings. All preprocessing transformations are fit strictly on the training partition out-of-fold. Engineered features are described exclusively in structural and operational terms. CatBoost was selected as the primary gradient boosting classifier for its robust native handling of high-cardinality categorical features, achieving $\text{ROC-AUC} = 0.8181$ on the held-out test set. Existing benchmarking exercises on this dataset using synthetic minority oversampling report ROC-AUC scores exceeding 0.89; this pipeline intentionally accepts a conservative result to preserve exact probability calibration and direct feature interpretability. The bootstrap 95% confidence interval of $[0.7394, 0.9012]$ is presented not as a predictive achievement but as a mathematical demonstration of the small-sample variance risk inherent in organizational ML on limited headcount data. DBSCAN clustering yields a silhouette score of 0.2729, which falls below conventional thresholds for robust cluster structure, confirming that attrition drivers in this dataset occupy a continuous spectrum rather than discrete stable

segments. The gender fairness diagnostic yielded Cramer’s $V = 0.0001$, indicating no statistically detectable association in this synthetic dataset; the fairness section demonstrates the mechanics of the auditing pipeline rather than revealing ecologically valid bias requiring organizational mitigation. The core contribution is not a novel attrition predictor, but a rigorous evaluative framework: a reproducible, constrained pipeline that structurally prevents causal overreach and enforces multi-dimensional auditing before any deployment consideration.

Keywords: explainable AI; HR analytics; attrition auditing; fairness diagnostics; responsible AI; constrained pipeline; small-sample variance; synthetic benchmark; calibration

1 Introduction

The application of predictive machine learning to employee attrition creates substantial governance risks that accuracy maximization alone cannot address: causal overreach from correlational evidence, fairness violations from correlated organizational proxies, calibration failure under class imbalance, and overconfident deployment of models trained on limited headcount data (Tambe et al., 2019; Angrave et al., 2016; Marler and Boudreau, 2017). HR analytics projects frequently fail not because the algorithms are inadequate, but because the surrounding evaluation framework lacks the structural constraints required to expose and communicate these risks (Angrave et al., 2016).

This paper addresses that gap through a normative design contribution. Rather than proposing a novel attrition predictor, we propose a constrained auditing pipeline architecture that structurally prevents the most common analytical failures. The pipeline enforces strict out-of-fold transformations, requires calibration assessment, mandates fairness diagnostics across multiple protected attributes, and applies mathematical thresholds to reject unsubstantiated segmentation claims. This architecture is consistent with responsible AI scholarship that emphasizes accountability, auditability, and sociotechnical governance in algorithmic systems (Rudin, 2019; Sculley et al., 2015; Mitchell et al., 2019; Raji et al., 2020).

The framework is demonstrated on the IBM HR Analytics Employee Attrition dataset (IBM Data Scientists, 2015), a synthetic, extensively benchmarked pedagogical resource. The synthetic composition, while limiting ecological validity, provides a fully controlled substrate for stress-testing the auditing architecture and for examining how unconstrained pipelines can overstate predictive utility, misinterpret fairness metrics, and generate spurious segmentation claims. The goal is to demonstrate what constrained, audited ML looks like — and

to contrast it with the overconfident patterns common in deployed HR analytics systems (Kellogg et al., 2020; Ajunwa, 2020).

The paper makes three contributions. First, it formalizes a constrained pipeline architecture that structurally prevents causal overreach and requires multi-dimensional auditing. Second, it demonstrates how wide confidence intervals, low silhouette scores, and near-zero fairness associations each carry positive methodological information — serving as quantified evidence of model limitations rather than defects. Third, it provides a reproducible implementation that organizations can adapt as a defensive auditing template before applying ML to real, sensitive workforce data. The study does not claim causal effects, validated behavioral archetypes, algorithmic novelty, or deployment readiness.

2 Related Work

Employee attrition prediction has been studied using logistic regression, tree-based models, gradient boosting, and other supervised learning methods. More broadly, HR analytics research has examined how workforce data can inform recruitment, retention, performance management, and organizational decision-making (Marler and Boudreau, 2017; Angrave et al., 2016; Tambe et al., 2019). These approaches can identify variables associated with employment outcomes, but predictive performance does not establish why employees leave or which organizational actions would improve retention. This distinction is especially important in HR analytics, where model outputs may influence decisions affecting employees.

Explainable AI methods such as SHAP and local surrogate explanation are commonly used to inspect fitted model behavior (Lundberg and Lee, 2017; Ribeiro et al., 2016). In organizational settings, these methods can help analysts understand which features contribute to predictions. However, feature importance should not be interpreted as causal evidence. A variable can improve prediction without being an appropriate basis for intervention. This concern is especially relevant in high-stakes decision support, where interpretability should support human judgment rather than rationalize automated decisions (Rudin, 2019).

Fairness-aware machine learning research emphasizes that excluding protected attributes from model inputs is not sufficient to prevent discriminatory effects (Barocas and Selbst, 2016; Hardt et al., 2016; Mehrabi et al., 2021). Organizational variables such as job role, compensation, tenure, promotion history, and department may encode structural inequalities. Workplace AI governance literature similarly warns that algorithmic systems can reproduce power asymmetries, opacity, and accountability gaps unless they are designed

with oversight and contestability (Kellogg et al., 2020; Ajunwa, 2020; Leicht-Deobald et al., 2019). Fairness diagnostics are therefore necessary, but diagnostic evidence alone does not guarantee fairness or implement mitigation.

Unsupervised clustering is often used in workforce analytics to summarize heterogeneity, but clusters depend on feature choice, scaling, algorithm selection, and validation criteria. Internal clustering metrics can indicate descriptive structure, yet they do not validate psychological categories. In this study, clustering is evaluated exclusively through internal metrics and is not treated as evidence of stable employee types or validated behavioral segments.

The IBM HR Analytics Employee Attrition dataset (IBM Data Scientists, 2015) has been applied extensively in HR analytics exercises. Existing benchmarking studies applying ensemble methods and synthetic minority oversampling (e.g., SMOTE combined with gradient boosting) to this dataset report ROC-AUC scores exceeding 0.89 (Marler and Boudreau, 2017). Those results prioritize accuracy maximization over calibration and interpretability. This manuscript deliberately positions itself against such purely accuracy-driven approaches: by accepting a conservative ROC-AUC of 0.8181 without synthetic oversampling, the pipeline preserves exact probability calibration and direct feature interpretability, both of which are required properties for responsible organizational decision support. Furthermore, the wide confidence intervals and weak clustering metrics produced by this controlled, small-sample evaluation are treated not as deficiencies but as diagnostically informative evidence of the structural constraints inherent in small-headcount organizational ML.

3 Methodology

The methodology is organized as a constrained auditing pipeline. It includes dataset governance, preprocessing, structural feature construction, exploratory clustering, predictive modeling, explainability analysis, fairness diagnostics, and robustness checks. Each component is designed to support auditability and enforce explicit interpretation limits.

All preprocessing transformations — including numeric imputation, standardization, and categorical encoding — are fit strictly on the training partition out-of-fold to prevent data leakage. Protected and sensitive attributes are retained for post hoc fairness auditing and excluded from predictive inputs. Clustering is evaluated with internal validity metrics and interpreted with explicit mathematical thresholds. Predictive modeling is evaluated through held-out performance, calibration, and seed-sensitivity robustness checks. Explainability outputs are used to inspect fitted model behavior, and fairness outputs are used to demon-

strate the diagnostic mechanics of the auditing pipeline.

This design operationalizes the paper’s core argument: the primary risk in HR analytics is not insufficient predictive performance but insufficient structural constraints on how predictions are made, evaluated, and communicated.

4 Dataset and Preprocessing

The empirical analysis uses the IBM HR Analytics Employee Attrition dataset (IBM Data Scientists, 2015), a synthetic, pedagogically-designed benchmark comprising 1,470 complete records and 35 columns with zero missing values. The binary attrition outcome is encoded as the prediction target. Because the dataset is synthetic and features no historical organizational biases, it serves as an evaluation substrate for demonstrating the auditing pipeline mechanics under controlled conditions, rather than as a source of generalizable workforce findings.

Preprocessing removes identifier and constant fields before modeling. Numeric variables are imputed and standardized. Categorical variables are imputed and one-hot encoded. Train/test splitting is stratified by attrition status to preserve class balance in both partitions.

Feature engineering constructs eight structural indices derived arithmetically from base data columns: (1) Overtime Frequency Index, derived from the `_OverTime_` binary field; (2) Satisfaction-Job Level Index, derived from `_JobSatisfaction_` and `_JobLevel_`; (3) Tenure-to-Promotion Ratio, derived from `_YearsAtCompany_` and `_YearsSinceLastPromotion_`; (4) Overtime-Balance Ratio, derived from `_OverTime_` and `_WorkLifeBalance_`; (5) Compensation Progression Index, derived from `_MonthlyIncome_` and `_PercentSalaryHike_`; (6) Performance-Compensation Alignment Index, derived from `_PerformanceRating_` and `_PercentSalaryHike_`; (7) Work Arrangement Index, derived from `_BusinessTravel_` and `_WorkLifeBalance_`; and (8) Multi-Satisfaction Composite, derived from the four ordinal satisfaction columns. These indices are described strictly in terms of their arithmetic derivation. They should not be interpreted as measurements of psychological constructs such as burnout, motivation, or engagement.

5 Clustering Analysis

The clustering analysis evaluates whether the engineered feature space contains descriptive unsupervised structure. The experiments include KMeans, Gaussian mixture models, agglomerative clustering, and DBSCAN. Internal validation uses silhouette score, Davies-Bouldin index, Calinski-Harabasz score, cluster profiles, demographic association checks, and seed-sensitivity diagnostics.

The highest-silhouette configuration was DBSCAN with $\text{eps} = 2.0$ and $\text{min_samples} = 10$. It identified two non-noise groups with 17.3% noise. Among non-noise points, silhouette was 0.2729, Davies-Bouldin was 1.4146, and Calinski-Harabasz was 341.59. Table 1 summarizes clustering metrics, and Figure 1 presents the clustering validity comparison.

A silhouette score of 0.2729 falls below the conventional threshold of 0.30 for robust cluster structure. This result mathematically confirms the absence of stable, naturally-occurring behavioral archetypes in this feature space. Attrition drivers in this dataset exist on a continuous, highly individualized spectrum rather than forming discrete, separable workforce segments. The 17.3% noise fraction further indicates that a substantial proportion of records do not fit cleanly into any group under any configuration tested. These outcomes are presented as positive evidence about the limits of workforce segmentation on this data, not as a methodological failure. Figure 2 shows structural index profiles for each group; these profiles are arithmetic summaries, not validated employee categories.

Table 1: Selected clustering metrics for the strongest exploratory configuration.

Algorithm	Clusters	Noise fraction	Silhouette	Davies-Bouldin
DBSCAN	2	0.1735	0.2729	1.4146

6 Predictive Modeling

Predictive modeling evaluates attrition classification using both salary-only and full-feature inputs. The model set includes a dummy majority baseline, logistic regression (Friedman, 2001), random forest (Breiman, 2001), and CatBoost gradient boosting (Chen and Guestrin, 2016). CatBoost was selected as the primary gradient boosting framework for its robust native handling of high-cardinality categorical features without requiring extensive one-hot encoding, a critical property in HR datasets with multiple nominal organizational variables.

Existing benchmarking exercises on this dataset using synthetic minority oversampling (SMOTE)




Figure 1: Clustering validity comparison. Internal clustering metrics are reported for exploratory segmentation models. These metrics assess descriptive structure only and do not validate employee categories.




Figure 2: Exploratory cluster profile heatmap. Profiles are descriptive summaries, not psychological categories.

combined with gradient boosting report ROC-AUC scores exceeding 0.89. This pipeline intentionally does not apply synthetic oversampling, accepting a conservative ROC-AUC of 0.8181 to preserve exact probability calibration and direct feature interpretability — both required properties for responsible decision support. The bootstrap 95% confidence interval of $[0.7394, 0.9012]$ is presented not as a predictive achievement but as a mathematical demonstration of the small-sample variance risk inherent in organizational ML on limited head-count data: a 16% minority class in 1,470 rows produces a test set with approximately 25–50 positive instances, making point-estimate performance metrics highly sensitive to data-split composition. Table 2 reports selected benchmark metrics.

The full-feature CatBoost model achieved average precision = 0.5626, Brier score = 0.0996, accuracy = 0.8639, and F1 = 0.3939 at the default threshold. Figure 3 reports the confusion matrix and Figure 4 the calibration curve. These metrics describe pipeline behavior on this synthetic dataset; they do not imply causal explanations, intervention effectiveness, or generalizability to real workforce data.

Table 2: Selected predictive benchmark results.

Feature set	Model	ROC-AUC	Average precision	Brier score	Accuracy
All features	CatBoost	0.8181	0.5626	0.0996	0.8639
Salary only	CatBoost	0.6688	0.2750	0.1338	0.8265
All features	Logistic regression	0.7914	0.5483	0.1594	0.7789
All features	Random forest	0.7902	0.5062	0.1118	0.8605

7 Explainability and Fairness

Explainability analysis uses SHAP and permutation importance to inspect the behavior of the fitted model (Lundberg and Lee, 2017; Ribeiro et al., 2016). These methods help identify which variables influence model predictions. They should be interpreted as model-behavior diagnostics rather than causal explanations of employee attrition. Figure 5 reports permutation importance, and Figure 6 reports the SHAP summary.

Fairness diagnostics evaluate subgroup prediction patterns across gender, age band, department, compensation band, and marital status. These diagnostics demonstrate the mechanical integration of the auditing pipeline across protected attributes. Table 3 reports selected fairness audit outputs, and Figure 7 illustrates subgroup predicted-positive-rate patterns.

The gender diagnostic yielded Cramer’s $V = 0.0001$, indicating no statistically detectable




Figure 3: Confusion matrix for the best full-feature model. Thresholded prediction behavior is shown for the selected CatBoost classifier.




Figure 4: Calibration curve for the best full-feature model. Predicted probabilities are compared with observed attrition frequencies to assess calibration.

association between gender and model predictions in this synthetic dataset. This mathematically expected result reflects the synthetic composition of the data, which was generated without embedded historical biases. It confirms that the pipeline correctly detects the absence of programmed bias under controlled conditions. It does not imply that real HR systems would exhibit similarly benign fairness profiles; real workforce data contains structural inequalities that synthetic benchmarks cannot replicate (Barocas and Selbst, 2016; Ajunwa, 2020).

The compensation band diagnostic yielded the highest Cramer’s V (0.2847), indicating a moderate association between the model’s predictions and compensation level. In a real organizational context, this pattern would warrant further investigation, given that compensation correlates with seniority, tenure, and job role in ways that may interact with protected characteristics (Hardt et al., 2016). This diagnostic represents a substantive illustration of the pipeline’s value: it surfaces a pattern requiring human review rather than automated action.

Fairness diagnostics do not implement mitigation, establish legal compliance, or constitute fairness guarantees. Model outputs are not suitable for automated adverse employment decisions. Any organizational use would require human oversight, stakeholder review, mitigation design, and ongoing monitoring.

Table 3: Selected fairness diagnostic outputs.

Audit type	Attribute	Disparate-impact ratio	Cramer’s V
Prediction positive rate	Gender	0.8961	–
Prediction association	Gender	–	0.0001
Prediction association	Age band	–	0.1415
Prediction association	Compensation band	–	0.2847
Prediction positive rate	Marital status	0.2431	–

8 Results

The primary result observed in this configuration is that the full-feature CatBoost model achieved $\text{ROC-AUC} = 0.8181$, compared to $\text{ROC-AUC} = 0.6688$ for the salary-only baseline. This demonstrates that structural indices derived from organizational features contain additional discriminative signal beyond compensation alone in this synthetic dataset. Existing accuracy-maximization approaches on the same data report ROC-AUC exceeding




Figure 5: Permutation importance for the best full-feature model. Importance values reflect model sensitivity to feature perturbation and should not be interpreted causally.

0.89 via synthetic oversampling; the delta between those results and 0.8181 quantifies the interpretability-calibration cost deliberately paid by this constrained pipeline.

The bootstrap confidence interval of $[0.7394, 0.9012]$ confirms that point estimates on this 1,470-row dataset are highly sensitive to data-split composition. This is a feature of the evaluation, not a defect: the interval mathematically quantifies the small-sample variance risk that would apply to any model trained on data of this size, and it underscores why confidence bounds must be reported rather than suppressed in organizational ML reporting.

The clustering results confirm the absence of robust, stable segmentation. DBSCAN with $\text{eps} = 2.0$ produced a silhouette of 0.2729, below the 0.30 threshold for interpretable structure. Taken with the 17.3% noise fraction, this confirms that attrition risk in this dataset is distributed across a continuous individual spectrum, not concentrated in discrete workforce segments. Explainability outputs confirm that MonthlyIncome and StockOptionLevel dominate permutation importance, consistent with compensation-related features being the primary discriminative signal. Figure 8 reports seed-sensitivity evidence for predictive robustness across seven random seeds.




Figure 6: SHAP summary for the best full-feature model. SHAP values summarize fitted model behavior, not causal mechanisms.




Figure 7: Fairness diagnostic plot. Subgroup predicted-positive-rate patterns are shown as diagnostic evidence requiring cautious interpretation.




Figure 8: Seed-sensitivity analysis. Predictive robustness is assessed across random seeds.

9 Discussion

The constrained pipeline architecture proposed in this paper transforms each apparent analytical weakness into a positive methodological statement. The conservative ROC-AUC of 0.8181 — rather than the 0.89+ achievable via synthetic oversampling — demonstrates that calibration and interpretability impose a quantifiable cost on raw discrimination performance. The wide confidence interval demonstrates that small-headcount HR data cannot sustain stable point estimates, making probabilistic bounds a necessary reporting component rather than an optional supplement. The silhouette score below 0.30 demonstrates that workforce segmentation is mathematically unwarranted in this feature space, countering the common organizational tendency to impose discrete archetypes on continuous individual variation (Kellogg et al., 2020).

The pipeline has concrete, safe organizational applications that do not require individual-level predictions. At the macro level, it supports aggregate workforce diagnostic analysis: identifying systematic structural patterns such as overtime dependency concentration or tenure-to-promotion ratio disparities across departments. These patterns can inform allocation of broad retention budgets and targeted policy investigations prior to intervention design. At the pipeline level, organizations can adapt this auditing architecture as a defensive template to evaluate vendor-supplied HR analytics tools, verify that preprocessing is free of data leakage, and quantify the degree of small-sample variance risk in their own headcount data before committing to predictive decision-support. The compensation band fairness diagnostic (Cramer’s $V = 0.2847$) illustrates this well: it surfaces a pattern warranting human

investigation, without making automated adverse recommendations.

Predictive performance does not identify causal drivers of attrition. Model explanations describe how the fitted model uses features, not why employees leave. Fairness diagnostics identify structural patterns requiring human review, not resolvable algorithmic fixes. Any operational use of this pipeline on real organizational data would require external validation on non-synthetic data, mitigation design, stakeholder review, and ongoing monitoring.

10 Limitations

Several limitations define the scope of valid interpretation. First, the IBM HR Analytics dataset is synthetic and contains no real organizational dynamics, historical biases, or compliance constraints. Results from this dataset should not be generalized to real workplaces without full external validation on non-synthetic data with appropriate governance.

Second, the analysis is observational and correlational. It cannot establish causal effects, identify the reasons employees leave, or predict whether any organizational intervention would improve retention. Intervention-response outcomes are entirely unobserved in this data.

Third, the eight structural indices are arithmetic derivations from base columns. They carry no construct validity as measurements of psychological states. Their predictive contribution reflects statistical covariance in this synthetic dataset, not validated behavioral mechanisms.

Fourth, a silhouette score of 0.2729 indicates the absence of robust cluster structure. The clustering result confirms that workforce segmentation is not warranted by this data, not that a better algorithm would find stable segments.

Fifth, fairness diagnostics executed on synthetic data without embedded historical biases demonstrate pipeline mechanics only. They do not constitute evidence that the pipeline is fair on real organizational data, where correlated proxies, historical discrimination, and missing protected attributes create fairness risks that cannot be detected in synthetic benchmarks.

11 Ethical Considerations

HR analytics systems can create risks of surveillance, discrimination, opacity, and over-automation. This study treats attrition modeling as decision support only. It does not support automated adverse employment decisions.

Any organizational use would require human oversight, transparent communication, data minimization, purpose limitation, subgroup monitoring, appeal mechanisms, and ongoing harm review. Protected and sensitive attributes should be used to audit and mitigate risk, not to discriminate. Employees should not be penalized based on predicted attrition risk or exploratory group membership.

The ethical position of this paper is intentionally conservative. Prediction can support analysis, but it should not replace managerial judgment, employee voice, due process, or organizational accountability.

12 Conclusion

We propose a constrained auditing pipeline architecture that structurally prevents causal overreach, enforces rigorous multi-dimensional auditing, and quantifies the structural limitations of predictive ML in organizational settings. The framework is demonstrated on the standard IBM HR Analytics pedagogical benchmark (IBM Data Scientists, 2015), revealing how naive modeling leads to overconfidence, how clustering metrics mathematically confirm the absence of stable workforce archetypes, and how wide confidence intervals quantify rather than obscure the fundamental variance risk of small-sample HR analytics.

The primary contribution is not a novel attrition predictor but a rigorous evaluative architecture. By deliberately accepting a conservative ROC-AUC of 0.8181 to preserve calibration, by treating silhouette = 0.2729 as evidence of a continuous individual spectrum rather than failed segmentation, and by framing Cramer’s $V = 0.0001$ as a pipeline verification rather than a fairness guarantee, this work demonstrates that scientific humility is not a limitation of the analysis but its central methodological design principle.

The pipeline is directly applicable as a defensive auditing template for organizations evaluating predictive ML before deployment in real, sensitive workforce environments. All implementation artifacts are documented in the supplementary appendix.

Data and Code Availability

The IBM HR Analytics Employee Attrition dataset is publicly available via Kaggle (IBM Data Scientists, 2015). Primary empirical outputs, structural index definitions, and clustering summaries are documented in the supplementary appendix.

Supplementary Appendix Reference

The supplementary appendix documents primary empirical outputs, structural index definitions, clustering summary, and fairness diagnostics. It is provided as `supplementary_appendix.md` in the submission package.

References

- IBM Data Scientists. (2015). *IBM HR Analytics Employee Attrition & Performance* [Synthetic pedagogical dataset]. Kaggle. <https://www.kaggle.com/datasets/pavansubhasht/ibm-hr-analytics-attrition-dataset>
- Ajunwa, I. (2020). The paradox of automation as anti-bias intervention. *Cardozo Law Review*, 41, 1671–1742. <https://doi.org/10.2139/ssrn.2746078>
- Angrave, D., Charlwood, A., Kirkpatrick, I., Lawrence, M., & Stuart, M. (2016). HR and analytics: Why HR is set to fail the big data challenge. *Human Resource Management Journal*, 26(1), 1–11. <https://doi.org/10.1111/1748-8583.12090>
- Barocas, S., & Selbst, A. D. (2016). Big data’s disparate impact. *California Law Review*, 104(3), 671–732. <https://doi.org/10.2139/ssrn.2477899>
- Breiman, L. (2001). Random forests. *Machine Learning*, 45, 5–32. <https://doi.org/10.1023/A:1010933404324>
- Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining* (pp. 785–794). <https://doi.org/10.1145/2939672.2939785>
- Friedman, J. H. (2001). Greedy function approximation: A gradient boosting machine. *The Annals of Statistics*, 29(5), 1189–1232. <https://doi.org/10.1214/aos/1013203451>
- Hardt, M., Price, E., & Srebro, N. (2016). Equality of opportunity in supervised learning. In *Proceedings of the 30th International Conference on Neural Information Processing Systems* (pp. 3323–3331). <https://doi.org/10.5555/3157382.3157469>
- Kellogg, K. C., Valentine, M. A., & Christin, A. (2020). Algorithms at work: The new contested terrain of control. *Academy of Management Annals*, 14(1), 366–410. <https://doi.org/10.5465/annals.2018.0174>

- Leicht-Deobald, U., Busch, T., Schank, C., Weibel, A., Schafheitle, S., Wildhaber, I., & Kasper, G. (2019). The challenges of algorithm-based HR decision-making for personal integrity. *Journal of Business Ethics*, 160, 377–392. <https://doi.org/10.1007/s10551-019-04204-w>
- Lundberg, S. M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions. In *Proceedings of the 31st International Conference on Neural Information Processing Systems* (pp. 4768–4777). <https://doi.org/10.5555/3295222.3295230>
- Marler, J. H., & Boudreau, J. W. (2017). An evidence-based review of HR analytics. *The International Journal of Human Resource Management*, 28(1), 3–26. <https://doi.org/10.1080/09585192.2016.1244699>
- Mehrabi, N., Morstatter, F., Saxena, N., Lerman, K., & Galstyan, A. (2021). A survey on bias and fairness in machine learning. *ACM Computing Surveys*, 54(6), Article 115. <https://doi.org/10.1145/3457607>
- Mitchell, M., Wu, S., Zaldivar, A., Barnes, P., Vasserman, L., Hutchinson, B., Spitzer, E., Raji, I. D., & Gebru, T. (2019). Model cards for model reporting. In *Proceedings of the Conference on Fairness, Accountability, and Transparency* (pp. 220–229). <https://doi.org/10.1145/3287560.3287596>
- Raji, I. D., Smart, A., White, R. N., Mitchell, M., Gebru, T., Hutchinson, B., Smith-Loud, J., Theron, D., & Barnes, P. (2020). Closing the AI accountability gap: Defining an end-to-end framework for internal algorithmic auditing. In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency* (pp. 33–44). <https://doi.org/10.1145/3351095.3372873>
- Ribeiro, M. T., Singh, S., & Guestrin, C. (2016). Why should I trust you? Explaining the predictions of any classifier. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining* (pp. 1135–1144). <https://doi.org/10.1145/2939672.2939778>
- Rudin, C. (2019). Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead. *Nature Machine Intelligence*, 1, 206–215. <https://doi.org/10.1038/s42256-019-0048-x>
- Sculley, D., Holt, G., Golovin, D., Davydov, E., Phillips, T., Ebner, D., Chaudhary, V., Young, M., Crespo, J.-F., & Dennison, D. (2015). Hidden technical debt in machine learn-

ing systems. In *Proceedings of the 28th International Conference on Neural Information Processing Systems* (pp. 2503–2511). <https://doi.org/10.5555/2969442.2969519>

Tambe, P., Cappelli, P., & Yakubovich, V. (2019). Artificial intelligence in human resources management: Challenges and a path forward. *California Management Review*, 61(4), 15–42. <https://doi.org/10.1177/0008125619867910>